

# Databases & SQL for Social Data

Jared Edgerton

**Note on data.** This week uses **synthetic** (simulated) campaign finance-style data generated in the course script. The goal is to practice database thinking and SQL—not to draw substantive inferences about real campaigns.

## Conceptual Questions

Please write three to ten sentence explanations for each of the following questions. **You are only required to answer ONE of the two questions below.**

1. Explain what a **relational schema** is and why it is useful for social data. In your answer, define **primary keys** and **foreign keys**, and explain how they reduce duplication and enable joins. Use the (candidate, contributor, contribution) setting from this week’s coding lab as your concrete example.
2. Explain what a database **index** does and why it can make queries faster. What are two trade-offs of adding indexes (e.g., disk usage, slower inserts/updates, maintenance)? Give one example of a query pattern (filter, join, group-by) that is likely to benefit from an index in the lab dataset.

## Applied Exercises

Use the code in the week’s code tutorial and the lecture slides to answer the following questions.

3. **Build the database + inspect the schema (synthetic data).** Using the provided @MastersThesis•, author = •, title = •, school = •, year = •, OPTkey = •, OPTtype = •, OPTaddress = •, OPTmonth = •, OPTnote = •, OPTannote = • script, create the SQLite database (`campaign.finance.db`) and load the synthetic tables: `candidates`, `contributors`, and `contributions`.
  - Report the row counts in each table using `SELECT COUNT(*)`.
  - Show the schema for each table (e.g., `PRAGMA table_info(candidates)` and similarly for the other two tables).
  - Briefly explain (2–4 sentences) how `contributor_id` and `candidate_id` connect the tables.
4. **Joins + aggregation (write and run your own SQL).** Write a SQL query (and run it through R or Python) that uses at least **one join** and at least **one aggregation**:
  - Required: join `contributions` to `candidates` and compute total contributions by `party`.
  - Required: restrict to contributions with `amount > 1000`.
  - Output: a clean table with columns `party`, `total_amount`, and `num_contributions`.
  - Visualization: make a simple bar plot of `total_amount` by `party`.

**5. Indexes + query plan (evidence of optimization).** Using SQL statements, do the following:

- Verify which indexes exist on `contributions` (e.g., query `sqlite_master`).
- Choose one query that filters by `candidate_id` or `date` or `amount`. Run EXPLAIN QUERY PLAN for that query.
- In 4–6 sentences, interpret the query plan: does SQLite report using an index? If not, what index might help and why?

**6. Challenge Question (Optional — if you finish early):** Run the same join-and-aggregate query in **DuckDB** (from R or Python) and compare it to SQLite.

- Repeat the query in DuckDB (you may load the tables from CSV or connect to the SQLite database, depending on what you set up in lab).
- Provide one piece of evidence about performance (e.g., a timing using `system.time()` in R or `timing` in Python).
- In 3–6 sentences, explain what factors might drive differences (data size, indexing, execution engine, I/O, query planner).